

US Accidents Exploratory Data Analysis

Description This is a countrywide car accident dataset that covers 49 states of the USA. The accident data were collected from February 2016 to March 2023, using multiple APIs that provide streaming traffic incident (or event) data. These APIs broadcast traffic data captured by various entities, including the US and state departments of transportation, law enforcement agencies, traffic cameras, and traffic sensors within the road networks. The dataset currently contains approximately 7.7 Million accident records. For more information about this dataset, please visit [here](#).

TODO - TALK ABOUT EDA TODO - TALK ABOUT DATASET (SOURCES,WHAT IT CONTAINS,HOW IT WILL BE USEFUL)

KAGGLE, INFORMATION ABOUT ACCIDENTS, CAN BE USEFUL TO PREVENT ACCIDENTS Mention that this does not contain data about New York

Importing the important python libraries

```
In [24]: # Importing the important python libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

DATA PREPARATION AND CLEANING

```
In [25]: df = pd.read_csv(r"C:\Users\AYUSH\Downloads\archive (1)\US_Accidents_March23.csv")
```

```
In [26]: df
```

```
Out [26]:
```

	ID	Source	Severity	Start_Time	End_Time	Start_Lat	Start_Lng	End_Lat	End_Lng	Distance(mi)	...	Roundabout	Station	Stop
0	A-1	Source2	3	2016-02-08 05:46:00	2016-02-08 11:00:00	39.865147	-84.058723	NaN	NaN	0.010	...	False	False	False
1	A-2	Source2	2	2016-02-08 06:07:59	2016-02-08 06:37:59	39.928059	-82.831184	NaN	NaN	0.010	...	False	False	False
2	A-3	Source2	2	2016-02-08 06:49:27	2016-02-08 07:19:27	39.063148	-84.032608	NaN	NaN	0.010	...	False	False	False
3	A-4	Source2	3	2016-02-08 07:23:34	2016-02-08 07:53:34	39.747753	-84.205582	NaN	NaN	0.010	...	False	False	False
4	A-5	Source2	2	2016-02-08 07:39:07	2016-02-08 08:09:07	39.627781	-84.188354	NaN	NaN	0.010	...	False	False	False
...
7728389	A-7777757	Source1	2	2019-08-23 18:03:25	2019-08-23 18:32:01	34.002480	-117.379360	33.99888	-117.37094	0.543	...	False	False	False
7728390	A-7777758	Source1	2	2019-08-23 19:11:30	2019-08-23 19:38:23	32.766960	-117.148060	32.76555	-117.15363	0.338	...	False	False	False
7728391	A-7777759	Source1	2	2019-08-23 19:00:21	2019-08-23 19:28:49	33.775450	-117.847790	33.77740	-117.85727	0.561	...	False	False	False
7728392	A-7777760	Source1	2	2019-08-23 19:00:21	2019-08-23 19:29:42	33.992460	-118.403020	33.98311	-118.39565	0.772	...	False	False	False
7728393	A-7777761	Source1	2	2019-08-23 18:52:06	2019-08-23 19:21:31	34.133930	-117.230920	34.13736	-117.23934	0.537	...	False	False	False

7728394 rows x 46 columns

```
In [27]: df.columns
```

```
Out [27]: Index(['ID', 'Source', 'Severity', 'Start_Time', 'End_Time', 'Start_Lat',
                'Start_Lng', 'End_Lat', 'End_Lng', 'Distance(mi)', 'Description',
                'Street', 'City', 'County', 'State', 'Zipcode', 'Country', 'Timezone',
                'Airport_Code', 'Weather_Timestamp', 'Temperature(F)', 'Wind_Chill(F)',
                'Humidity(%)', 'Pressure(in)', 'Visibility(mi)', 'Wind_Direction',
                'Wind_Speed(mph)', 'Precipitation(in)', 'Weather_Condition', 'Amenity',
                'Bump', 'Crossing', 'Give_Way', 'Junction', 'No_Exit', 'Railway',
                'Roundabout', 'Station', 'Stop', 'Traffic_Calming', 'Traffic_Signal',
                'Turning_Loop', 'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight',
                'Astronomical_Twilight'],
                dtype='object')
```

In [28]: # Overall Information about the Data

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7728394 entries, 0 to 7728393
Data columns (total 46 columns):
#   Column              Dtype
---  -
0   ID                   object
1   Source               object
2   Severity             int64
3   Start_Time           object
4   End_Time             object
5   Start_Lat            float64
6   Start_Lng            float64
7   End_Lat              float64
8   End_Lng              float64
9   Distance(mi)         float64
10  Description           object
11  Street               object
12  City                 object
13  County               object
14  State                object
15  Zipcode              object
16  Country              object
17  Timezone             object
18  Airport_Code         object
19  Weather_Timestamp    object
20  Temperature(F)       float64
21  Wind_Chill(F)        float64
22  Humidity(%)          float64
23  Pressure(in)         float64
24  Visibility(mi)       float64
25  Wind_Direction       object
26  Wind_Speed(mph)      float64
27  Precipitation(in)    float64
28  Weather_Condition    object
29  Amenity              bool
30  Bump                 bool
31  Crossing             bool
32  Give_Way             bool
33  Junction             bool
34  No_Exit              bool
35  Railway              bool
36  Roundabout          bool
37  Station              bool
38  Stop                 bool
39  Traffic_Calming      bool
40  Traffic_Signal       bool
41  Turning_Loop         bool
42  Sunrise_Sunset       object
43  Civil_Twilight       object
44  Nautical_Twilight    object
45  Astronomical_Twilight object
dtypes: bool(13), float64(12), int64(1), object(20)
memory usage: 2.0+ GB
```

In [29]: # Descriptive Statistics of all the numerical columns

```
df.describe()
```

Out [29]:

	Severity	Start_Lat	Start_Lng	End_Lat	End_Lng	Distance(mi)	Temperature(F)	Wind_Chill(F)	Humidity(%)	Pressure
count	7.728394e+06	7.728394e+06	7.728394e+06	4.325632e+06	4.325632e+06	7.728394e+06	7.564541e+06	5.729375e+06	7.554250e+06	7.587715e+06
mean	2.212384e+00	3.620119e+01	-9.470255e+01	3.626183e+01	-9.572557e+01	5.618423e-01	6.166329e+01	5.825105e+01	6.483104e+01	2.953899e+01
std	4.875313e-01	5.076079e+00	1.739176e+01	5.272905e+00	1.810793e+01	1.776811e+00	1.901365e+01	2.238983e+01	2.282097e+01	1.006190e+01
min	1.000000e+00	2.455480e+01	-1.246238e+02	2.456601e+01	-1.245457e+02	0.000000e+00	-8.900000e+01	-8.900000e+01	1.000000e+00	0.000000e+00
25%	2.000000e+00	3.339963e+01	-1.172194e+02	3.346207e+01	-1.177543e+02	0.000000e+00	4.900000e+01	4.300000e+01	4.800000e+01	2.937000e+01
50%	2.000000e+00	3.582397e+01	-8.776662e+01	3.618349e+01	-8.802789e+01	3.000000e-02	6.400000e+01	6.200000e+01	6.700000e+01	2.986000e+01
75%	2.000000e+00	4.008496e+01	-8.035368e+01	4.017892e+01	-8.024709e+01	4.640000e-01	7.600000e+01	7.500000e+01	8.400000e+01	3.003000e+01
max	4.000000e+00	4.900220e+01	-6.711317e+01	4.907500e+01	-6.710924e+01	4.417500e+02	2.070000e+02	2.070000e+02	1.000000e+02	5.863000e+01

In [30]:

```
numerics = ['int16', 'int32', 'int64', 'float16', 'float32', 'float64']
numeric_df = df.select_dtypes(include=numerics)
len(numeric_df.columns)
```

Out [30]: 13

In [31]: # Missing entry

```
df.isna().sum()
```

Out [31]:

ID	0
Source	0
Severity	0
Start_Time	0
End_Time	0
Start_Lat	0
Start_Lng	0
End_Lat	3402762
End_Lng	3402762
Distance(mi)	0
Description	5
Street	10869
City	253
County	0

```

State      1915
Zipcode    0
Country    0
Timezone   7808
Airport_Code 22635
Weather_Timestamp 120228
Temperature(F) 163853
Wind_Chill(F) 1999019
Humidity(%) 174144
Pressure(in) 140679
Visibility(mi) 177098
Wind_Direction 175206
Wind_Speed(mph) 571233
Precipitation(in) 2203586
Weather_Condition 173459
Amenity    0
Bump       0
Crossing   0
Give_Way   0
Junction   0
No_Exit    0
Railway    0
Roundabout 0
Station    0
Stop       0
Traffic_Calming 0
Traffic_Signal 0
Turning_Loop 0
Sunrise_Sunset 23246
Civil_Twilight 23246
Nautical_Twilight 23246
Astronomical_Twilight 23246
dtype: int64

```

```

In [32]: #percentage of missing values per columns
missing_df = df.isna().sum().sort_values(ascending=False)/len(df)

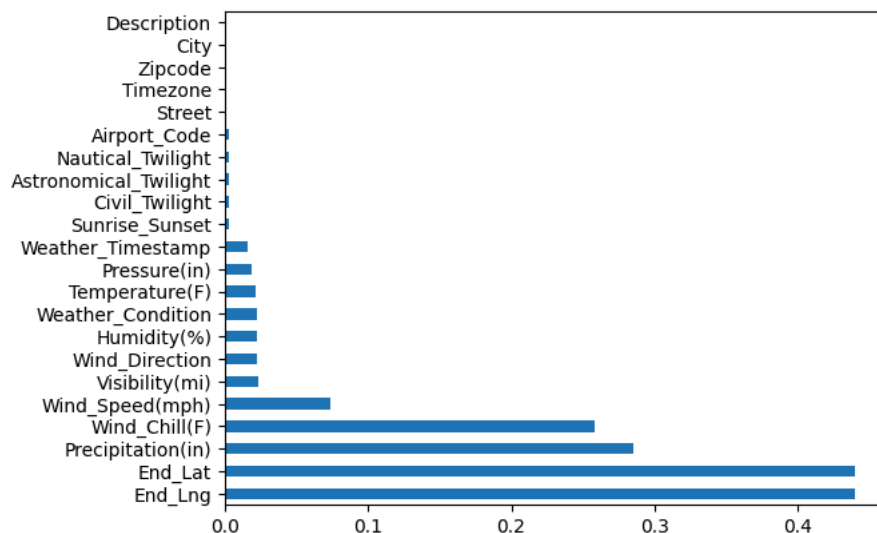
```

```

In [33]: missing_df[missing_df!=0].plot(kind='barh')

```

Out [33]: <Axes: >



```

In [34]: percentage_miss = missing_df[missing_df != 0]
percentage_miss

```

```

Out [34]: End_Lng      4.402935e-01
End_Lat      4.402935e-01
Precipitation(in) 2.851286e-01
Wind_Chill(F) 2.586590e-01
Wind_Speed(mph) 7.391355e-02
Visibility(mi) 2.291524e-02
Wind_Direction 2.267043e-02
Humidity(%) 2.253301e-02
Weather_Condition 2.244438e-02
Temperature(F) 2.120143e-02
Pressure(in) 1.820288e-02
Weather_Timestamp 1.555666e-02
Sunrise_Sunset 3.007869e-03
Civil_Twilight 3.007869e-03
Astronomical_Twilight 3.007869e-03
Nautical_Twilight 3.007869e-03
Airport_Code 2.928810e-03
Street 1.406372e-03
Timezone 1.010300e-03
Zipcode 2.477876e-04
City 3.273643e-05
Description 6.469649e-07
dtype: float64

```

remove columns that are not useful

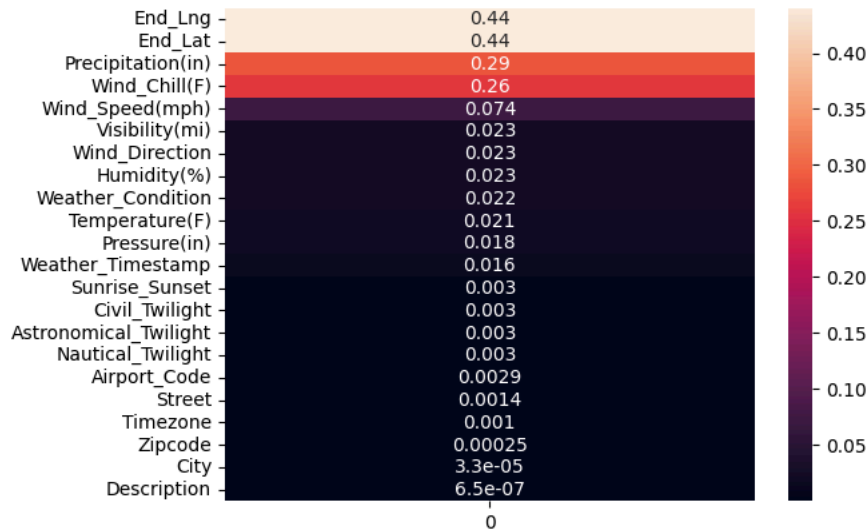
```

In [35]: # Converting Series to Dataframe
percentage_miss = percentage_miss.to_frame()

```

```
In [36]: # Heatmap Visualization
sns.heatmap(percentage_miss,annot=True)
```

Out [36]: <Axes: >



```
In [37]: # # We will exclude the columns that contain more than 5% of missing values, as they are not app

cols = [var for var in df.columns if df[var].isnull().mean(>0.05)]
cols
```

Out [37]: ['End_Lat', 'End_Lng', 'Wind_Chill(F)', 'Wind_Speed(mph)', 'Precipitation(in)']

```
In [38]: # Unnecessary columns removal
df.drop(columns = cols,inplace = True)
```

Exploratory Data And Visualisation

columns we will analyse

1. city
2. start time
3. start lat,start lng
4. state
5. temperature
6. weather condition

```
In [39]: df.columns
```

Out [39]: Index(['ID', 'Source', 'Severity', 'Start_Time', 'End_Time', 'Start_Lat', 'Start_Lng', 'Distance(mi)', 'Description', 'Street', 'City', 'County', 'State', 'Zipcode', 'Country', 'Timezone', 'Airport_Code', 'Weather_Timestamp', 'Temperature(F)', 'Humidity(%)', 'Pressure(in)', 'Visibility(mi)', 'Wind_Direction', 'Weather_Condition', 'Amenity', 'Bump', 'Crossing', 'Give_Way', 'Junction', 'No_Exit', 'Railway', 'Roundabout', 'Station', 'Stop', 'Traffic_Calming', 'Traffic_Signal', 'Turning_Loop', 'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight', 'Astronomical_Twilight'], dtype='object')

City

```
In [40]: cities = df.City.nunique()
cities
```

Out [40]: 13678

```
In [41]: cities_by_accident = df.City.value_counts()
cities_by_accident
```

Out [41]:

City	
Miami	186917
Houston	169609
Los Angeles	156491
Charlotte	138652
Dallas	130939
...	
Rapid River	1
Cat Spring	1
Glenwood City	1
Downing	1
Marfa	1
Name: count, Length: 13678, dtype: int64	

```
In [42]: cities_by_accident[:10]
```

```
Out [42]: City
Miami      186917
Houston     169609
Los Angeles 156491
Charlotte  138652
Dallas      130939
Orlando     109733
Austin       97359
Raleigh     86079
Nashville   72930
Baton Rouge 71588
Name: count, dtype: int64
```

```
In [43]: "New York" in df.City
```


```
Out [43]: False
```

```
In [44]: "NY" in df.State
```

```
Out [44]: False
```

```
In [45]: # Plot
cities_by_accident[:20].plot(kind="barh",title='Top 20 cities having most number of accidents',x
```

```
Out [45]: <Axes: title={'center': 'Top 20 cities having most number of accidents'}, xlabel='Number of Accidents',
ylabel='Cities'>
```



```
In [46]: sns.set_style("darkgrid")
```

```
In [47]: sns.distplot(cities_by_accident)
```

C:\Users\AYUSH\AppData\Local\Temp\ipykernel_54680\3405282844.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(cities_by_accident)
```

```
Out [47]: <Axes: xlabel='count', ylabel='Density'>
```




```
In [48]: high_accident_cities = cities_by_accident[cities_by_accident >= 1000]
low_accident_cities = cities_by_accident[cities_by_accident < 1000]
```

```
In [49]: len(high_accident_cities)/len(cities_by_accident)
```

Out [49]: 0.08904810644831115

```
In [51]: # Distribution II(Accidents are high)
sns.histplot(high_accident_cities,bins=100,log_scale=True,kde=True,.set(xlabel='Number of Accid
```

Out [51]: [Text(0.5, 0, 'Number of Accidents'), Text(0, 0.5, 'Number of Cities')]



```
In [52]: cities_by_accident[cities_by_accident==1]
```

Out [52]:

City	
American Fork-Pleasant Grove	1
Waldoboro	1
Kinsley	1
Killona	1
Jeanerette	1
..	
Rapid River	1
Cat Spring	1
Glenwood City	1
Downing	1
Marfa	1

Name: count, Length: 1023, dtype: int64

```
In [53]: sns.distplot(low_accident_cities)
```

C:\Users\AYUSH\AppData\Local\Temp\ipykernel_54680\469555131.py:1: UserWarning:


'distplot' is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either 'displot' (a figure-level function with similar flexibility) or 'histplot' (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

sns.distplot(low_accident_cities)

Out [53]: <Axes: xlabel='count', ylabel='Density'>




```
In [54]: # Number of Accidents reported in New York
(df[df['City']=='New York'].shape[0])
```

Out [54]: 21699

```
In [58]: # Distribution I(Accidents-City Analysis)
sns.set_style('darkgrid')
sns.histplot(cities_by_accident,bins=100,log_scale=True,kde=True,).set(xlabel='Number of Acciden
```

Out [58]: [Text(0.5, 0, 'Number of Accidents'), Text(0, 0.5, 'Number of Cities')]



Start Time

```
In [59]: df.columns
```

Out [59]: Index(['ID', 'Source', 'Severity', 'Start_Time', 'End_Time', 'Start_Lat', 'Start_Lng', 'Distance(mi)', 'Description', 'Street', 'City', 'County', 'State', 'Zipcode', 'Country', 'Timezone', 'Airport_Code', 'Weather_Timestamp', 'Temperature(F)', 'Humidity(%)', 'Pressure(in)', 'Visibility(mi)', 'Wind_Direction', 'Weather_Condition', 'Amenity', 'Bump', 'Crossing', 'Give_Way', 'Junction', 'No_Exit', 'Railway', 'Roundabout', 'Station', 'Stop', 'Traffic_Calming', 'Traffic_Signal', 'Turning_Loop', 'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight', 'Astronomical_Twilight'], dtype='object')


```
In [60]: df.Start_Time[0]
```

Out [60]: '2016-02-08 05:46:00'

```
In [61]: # Datatype Conversion
time=df['Start_Time']
time = pd.to_datetime(time, errors='coerce')
```

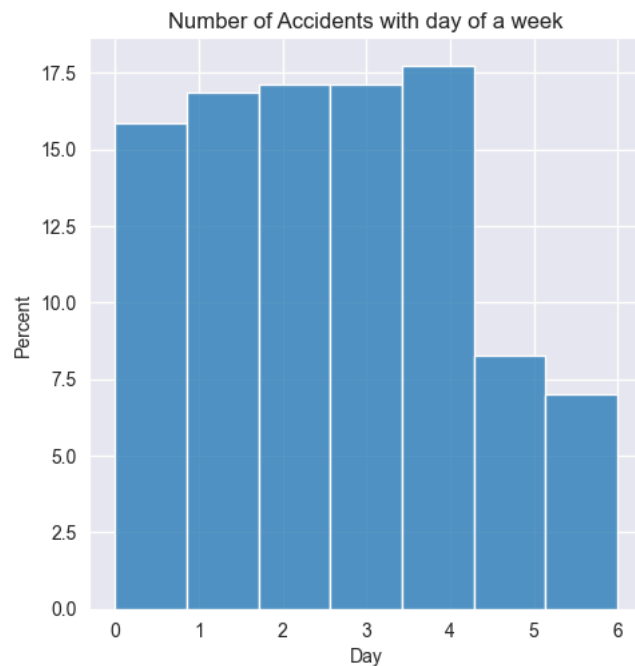
```
In [62]: # # Number of Accidents with hour of a day
sns.displot(time.dt.hour, bins=24,stat='percent')
plt.title('Number of Accidents with hour of a day')
```

```
plt.xlabel('Hour')
plt.show()
```



High Percentage Of Accidents occur between 6 to 10 am.(Probably people are in a hurry to get to work) Next Highest percentage is 3 to 6 pm.

```
In [63]: # Number of Accidents with day of a week
sns.displot(time.dt.dayofweek, bins=7, stat='percent')
plt.title('Number of Accidents with day of a week')
plt.xlabel('Day')
plt.show()
```



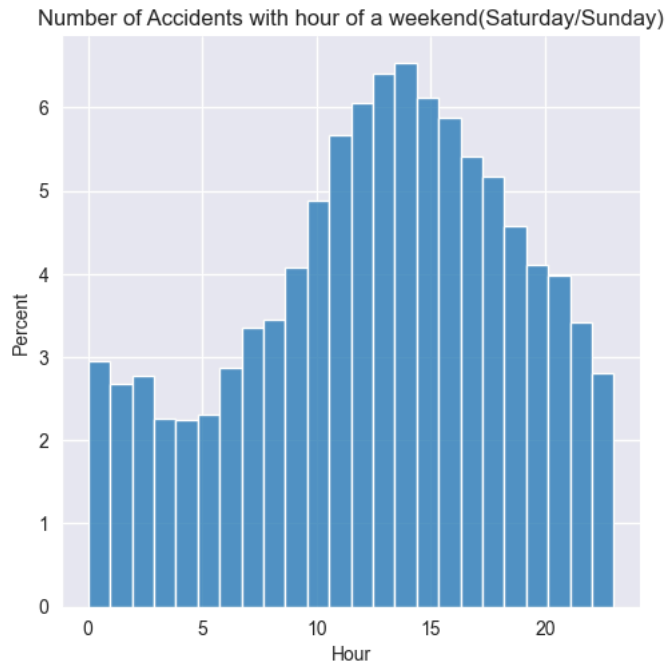
Is the distribution of accidents by hour the same on weekends as on weekdays.

```
In [64]: # Weekends-Accident Analysis
weekends = time[(time.dt.dayofweek==6) | (time.dt.dayofweek==5)]
weekends
```

```
Out [64]: 129      2016-02-13 11:05:00
130      2016-02-13 11:05:21
131      2016-02-13 11:17:01
132      2016-02-13 11:25:42
133      2016-02-13 12:56:31
...
7726137   2019-08-18 23:24:10
7726177   2019-08-17 03:36:35
7726252   2019-08-18 22:56:56
7726253   2019-08-18 22:56:56
7726292   2019-08-18 22:54:41
Name: Start_Time, Length: 1069268, dtype: datetime64[ns]
```

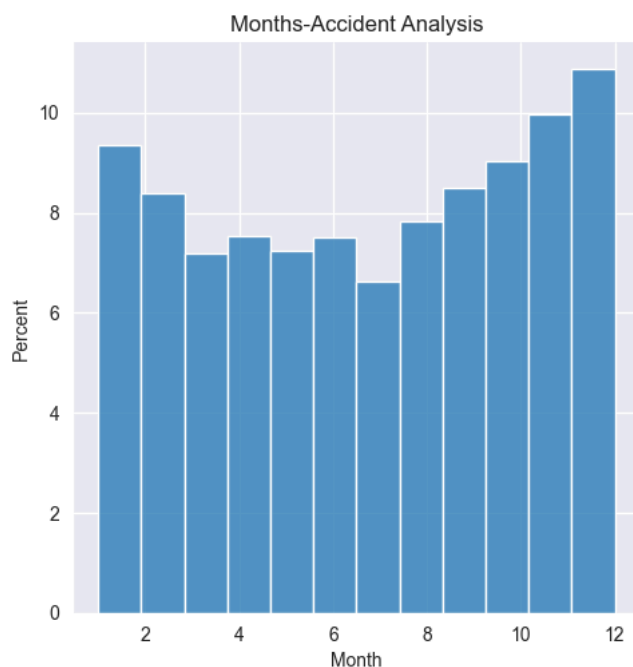


```
In [65]: # Number of Accidents with hour of a weekend(Saturday/Sunday)
sns.displot(weekends.dt.hour, bins=24,stat='percent')
plt.title('Number of Accidents with hour of a weekend(Saturday/Sunday)')
plt.xlabel('Hour')
plt.show()
```




On weekends, peak occur between 10 am to 3 pm, unlike weekdays.

```
In [66]: # Months-Accident Analysis
sns.displot(time.dt.month, bins=12,stat='percent')
plt.title('Months-Accident Analysis')
plt.xlabel('Month')
plt.show()
```



The majority of accidents tend to occur at the end of the year.

```
In [67]: # Years-Accident Analysis
sns.displot(time.dt.year, stat='percent')
plt.title('Years-Accident Analysis')
plt.xlabel('Year')
plt.show()
```



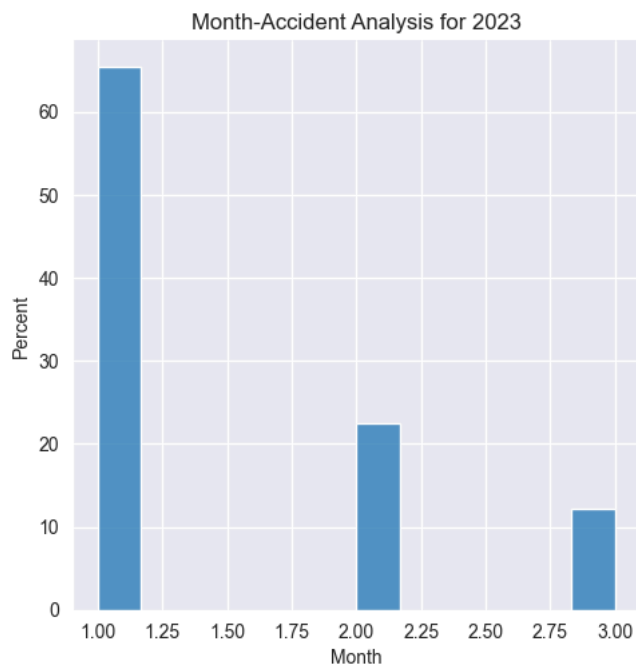
The number of accidents has been increasing each year, yet 2023 recorded the lowest number of accidents. This is highly unusual. We need to examine the data for 2023.

```
In [68]: # 2023 Analysis
```

```
time_2023 = time[time.dt.year==2023]
```

```
In [69]: # Visualization Plot
```

```
sns.displot(time_2023.dt.month,bins = 12,stat='percent')
plt.title('Month-Accident Analysis for 2023')
plt.xlabel('Month')
plt.show()
```



The reason for the seemingly lower number of accidents in 2023 is primarily due to data availability


```
In [70]: df.Source
```

```
Out [70]: 0      Source2
          1      Source2
          2      Source2
          3      Source2
          4      Source2
          ...
          7728389 Source1
          7728390 Source1
          7728391 Source1
          7728392 Source1
```

7728393 Source1
Name: Source, Length: 7728394, dtype: object

```
In [71]: df.Source.value_counts().plot(kind='pie')
```

Out [71]: <Axes: ylabel='count'>



Start Latitude & Longitude

```
In [72]: df.Start_Lng
```

Out [72]:

0	-84.058723
1	-82.831184
2	-84.032608
3	-84.205582
4	-84.188354
...	
7728389	-117.379360
7728390	-117.148060
7728391	-117.847790
7728392	-118.403020
7728393	-117.230920

Name: Start_Lng, Length: 7728394, dtype: float64

```
In [73]: df.Start_Lat
```

Out [73]:

0	39.865147
1	39.928059
2	39.063148
3	39.747753
4	39.627781
...	
7728389	34.002480
7728390	32.766960
7728391	33.775450
7728392	33.992460
7728393	34.133930

Name: Start_Lat, Length: 7728394, dtype: float64

```
In [74]: # Latitude
lat = df['Start_Lat']
lat
```

Out [74]:

0	39.865147
1	39.928059
2	39.063148
3	39.747753
4	39.627781
...	
7728389	34.002480
7728390	32.766960
7728391	33.775450
7728392	33.992460
7728393	34.133930

Name: Start_Lat, Length: 7728394, dtype: float64

```
In [75]: # Longitude
lon = df['Start_Lng']
lon
```


Out [75]:

0	-84.058723
1	-82.831184
2	-84.032608
3	-84.205582
4	-84.188354
...	
7728389	-117.379360
7728390	-117.148060
7728391	-117.847790
7728392	-118.403020
7728393	-117.230920

Name: Start_Lng, Length: 7728394, dtype: float64

```
In [76]: # Approximate Map of US
sns.scatterplot(x=df.Start_Lng,y=df.Start_Lat)
```

Out [76]: <Axes: xlabel='Start_Lng', ylabel='Start_Lat'>




```
In [78]: # Using folium to visualize the location of accidents on the map
import folium

from folium.plugins import HeatMap
lat_lon_pairs = list(zip(list(lat), list(lon)))

map = folium.Map()
HeatMap(lat_lon_pairs).add_to(map)
map
```

Out [78]:



State

```
In [79]: state = df['State']
```

```
In [80]: state.nunique()
```

Out [80]: 49


```
In [81]: # The top 5 states with the highest number of accidents
top_state = state.value_counts().head(5).reset_index()
```

```
In [82]: # VisualizationPlot
plt.figure(figsize=(8, 6))
sns.barplot(x='State', y='count', data=top_state, palette="magma")
plt.title('Top 5 highest Accident States')
plt.xlabel('State')
plt.ylabel('Count')
plt.show()
```

C:\Users\AYUSH\AppData\Local\Temp\ipykernel_54680\941678388.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='State', y='count', data=top_state, palette="magma")
```



Temperature

```
In [83]: df['Temperature(F)'].describe()
```

```
Out [83]: count    7.564541e+06
mean      6.166329e+01
std       1.901365e+01
min      -8.900000e+01
25%       4.900000e+01
50%       6.400000e+01
75%       7.600000e+01
max       2.070000e+02
Name: Temperature(F), dtype: float64
```

```
In [84]: # Checking the maximum recorded temperature
df[df['Temperature(F)']==207]
```

```
Out [84]:
```

	ID	Source	Severity	Start_Time	End_Time	Start_Lat	Start_Lng	Distance(mi)	Description	Street	...	Roundabout
4208142	A-4239885	Source1	2	2023-02-22 19:24:00.000000000	2023-02-22 21:32:35.000000000	34.066720	-117.238034	0.021	Accident on I-10 W from California St (I-10) t...	I-10 W	...	False
4546366	A-4580836	Source1	2	2023-02-22 20:10:00.0000000	2023-02-23 00:01:38.0000000	34.228800	-117.301024	0.417	Accident on CA-18 from N Sierra Way (N Waterma...	Rim of the World Hwy	...	False
5029338	A-5067455	Source1	4	2023-02-22 19:30:00	2023-02-23 00:00:53	34.228728	-117.251234	0.543	CA-189 is closed from CA-18/Lake Gregory Dr (C...	Lake Gregory Dr	...	False

3 rows x 41 columns

```
In [85]: temperature = df['Temperature(F)'].value_counts().reset_index()
```

```
In [86]: temperature = temperature.round(2)
temperature
```

```
Out [86]:
```

	Temperature(F)	count
0	77.0	170991
1	73.0	170898
2	68.0	163767
3	72.0	160498
4	75.0	158448
...
855	-32.8	1
856	-9.2	1
857	-17.9	1
858	168.8	1
859	-12.1	1

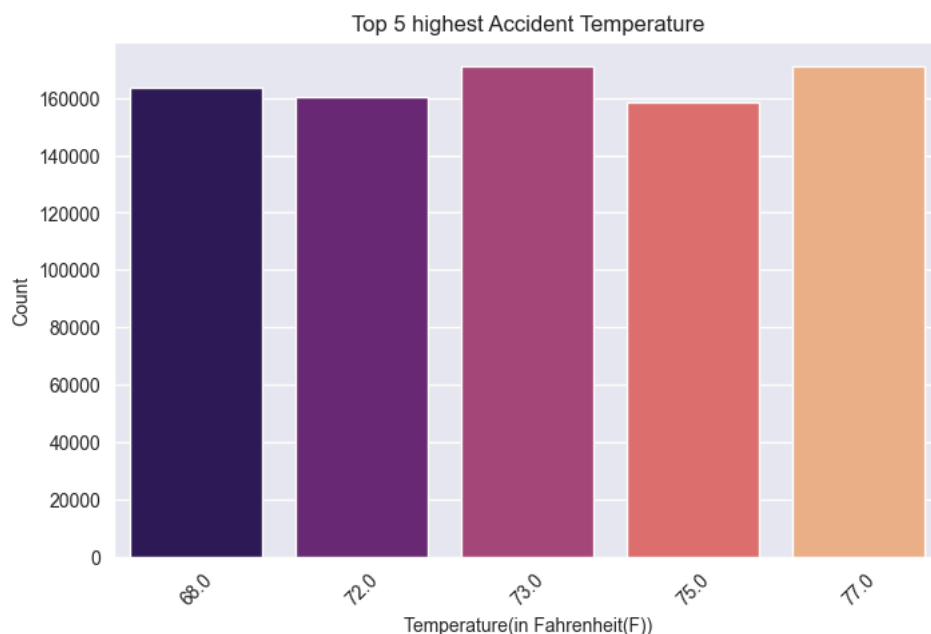
860 rows × 2 columns

```
In [87]: plt.figure(figsize=(8, 5))
sns.barplot(x='Temperature(F)', y='count', data=temperature.head(5), palette="magma")
plt.title('Top 5 highest Accident Temperature')
plt.xlabel('Temperature(in Fahrenheit(F))')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='center')
plt.show()
```

C:\Users\AYUSH\AppData\Local\Temp\ipykernel_54680\3890819157.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Temperature(F)', y='count', data=temperature.head(5), palette="magma")
```



```
In [88]: # Checking for data inconsistencies
temperature[temperature['Temperature(F)'] > 140].sum()
```

```
Out [88]: Temperature(F)    2272.6
count                22.0
dtype: float64
```

Weather

```
In [89]: df['Weather_Condition']
```

```
Out [89]: 0    Light Rain
1    Light Rain
2    Overcast
3    Mostly Cloudy
4    Mostly Cloudy
```

```

7728389      Fair
7728390      Fair
7728391  Partly Cloudy
7728392      Fair
7728393      Fair
Name: Weather_Condition, Length: 7728394, dtype: object

```

```

In [90]: weather = df['Weather_Condition'].value_counts().reset_index()

```

```

In [91]: # Top 10 Weather Conditions with most number of Accidents
weather = weather.head(10)
weather

```

```

Out [91]:   Weather_Condition  count
0      Fair           2560802
1  Mostly Cloudy       1016195
2    Cloudy           817082
3    Clear           808743
4  Partly Cloudy       698972
5  Overcast          382866
6  Light Rain         352957
7  Scattered Clouds   204829
8  Light Snow         128680
9    Fog              99238

```

```


In [92]: # Visualization Plot
plt.figure(figsize=(8, 6))
sns.barplot(x='Weather_Condition', y='count', data=weather, palette="magma")
plt.title('Top 5 highest Accident Weather Conditions')
plt.xlabel('State')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right')
plt.show()

```

C:\Users\AYUSH\AppData\Local\Temp\ipykernel_54680\2376085648.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Weather_Condition', y='count', data=weather, palette="magma")
```



Impact of Weather on Accident Severity:


```
In [93]: # Creating a DataFrame that summarizes the mean severity for each weather condition by grouping.
severity_weather = df.groupby('Weather_Condition')['Severity'].mean().reset_index().sort_values(
```

```
In [94]: # Visualisation of the Top 10 Weather Conditions that are associated with severe accidents.
plt.figure(figsize=(8, 6))
sns.barplot(x='Severity', y='Weather_Condition', data=severity_weather[0:10], palette="magma")
plt.title('Top 10 Weather Condition associated with severe accidents')
plt.xlabel('Mean Severity')
plt.ylabel('Weather Condition')
plt.show()
```

C:\Users\AYUSH\AppData\Local\Temp\ipykernel_54680\3304834884.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Severity', y='Weather_Condition', data=severity_weather[0:10], palette="magma")
```



Ask & Answer Questions

1. are there more accidents in warmer and cooler areas?
2. which 5 states have the highest numbers of accidents? how about per capita?
3. does New York show up in the data? if yes, why the count lower if this the most populated city?
4. Among top 100 cities in number of accidents, which states do they belong the most frequently?
5. What time of the day are accidents most frequent in ?
6. which day of the week have the most accidents?
7. which months have the most accidents?
8. what is the trends of accidents over the years (increase/decrease)?
9. when is accidents per unit of traffic the highest.

Summary & Conclusions

1. Approximately 9% of the cities experience more than 1,000 accidents.
2. About 1,023 cities (7%) have reported only one accident, suggesting the presence of potential outliers that may need to be addressed.
3. Miami leads with the highest number of accidents among the top 20 cities, while Columbia has the fewest accidents in this group.
4. New York accounts for only 21,699 accidents, which is approximately 0.28% of the total number of accidents in the US, despite being the most populated city in the country.
5. California, Florida, Texas, South Carolina, and New York emerge as the top 5 states with the highest number of accidents.
6. A high percentage of accidents occur between between 6 AM to 10 AM.
7. On Weekends, high percentage of accidents occur between 12PM to 4 PM with the peak accidents around 2PM.
8. Weekdays show a higher number of accidents compared to weekends.
9. The majority of accidents tend to occur at the end of the year.
10. The dataset used for analysis only includes information up to March 2023, which means it does not encompass the entire year. Consequently, the observed decrease in accidents for 2023 may not accurately represent the true accident rate for the entire year, as data for the subsequent months is missing. Therefore, any conclusions about accident trends in 2023 should be approached with caution, given the incomplete data for that year.

11. The top 10 most common weather conditions during accidents include: Fair, Mostly Cloudy, Cloudy, Clear, Partly Cloudy, Overcast, Light Rain, Scattered Clouds, Light Snow, and Fog.
12. Light Blowing Snow exhibits the highest average severity at 3.67.
13. Partial Fog with Wind, Heavy Thunderstorms and Snow, Light Thunderstorms and Snow, Heavy Ice Pellets, Heavy Blowing Snow, and Drifting Snow with Wind also demonstrate notable average severity levels.
14. Some weather conditions can significantly reduce visibility and create hazardous driving conditions. Heavy Thunderstorms and Snow are characterized by the lowest visibility among all weather conditions.
15. The initial analysis suggests a potential correlation between temperature and accident frequency within the range of 61.66°F and 77°F.
16. The dataset does contain approximately 22 entries that appear to be outliers based on expected temperature ranges.(The maximum reported temperature of US is around 134°F)
17. There is a seasonal variation in accidents, with fewer incidents during the summer and an increasing trend as winter approaches.
18. Accident rates are higher in coastal/bay areas than in inland regions.

Recommendations:

1. Safety Campaigns
 - Focus on cities with over 1,000 accidents.
 - Enhance awareness and enforcement during peak times (12 PM-4 PM weekends, 6 AM-10 AM weekdays).
2. Outlier Review
 - Investigate the 1,023 cities with only one accident for data accuracy.
3. Urban Traffic Management
 - Monitor and improve traffic strategies in major cities like New York.
4. State-Specific Programs
 - Develop safety programs for California, Florida, Texas, South Carolina, and New York.
5. Weather-Related Measures
 - Improve road safety and visibility during severe weather conditions.
 - Enhance road lighting in fog-prone areas.
6. Seasonal Adjustments
 - Increase safety measures for winter; analyze factors for fewer summer accidents.
7. Temperature Analysis
 - Investigate temperature-related accident trends and address outlier data.
8. Coastal Area Safety
 - Implement additional safety measures in coastal and bay areas.
9. Data Completeness
 - Ensure future datasets include full-year data for accurate trend analysis.